

AI Full-Stack Engineer interview preparation

# DocMind **AI** Interview Study Guide One-Day Edition

The project, the decisions, the essential concepts and the most likely questions — everything you need, readable in a day.

Built from the DocMind AI repository · 29 September 2026

# Contents

---

|                                                               |           |
|---------------------------------------------------------------|-----------|
| <b>1. The one-day plan and your pitch .....</b>               | <b>3</b>  |
| How to spend the day .....                                    | 3         |
| The 60-second pitch .....                                     | 3         |
| The 2-minute extension (if they say "tell me more") .....     | 3         |
| Numbers to know .....                                         | 4         |
| <b>2. How DocMind works .....</b>                             | <b>5</b>  |
| Upload → ingestion .....                                      | 6         |
| Question → streamed answer .....                              | 7         |
| Data model (5 tables) .....                                   | 8         |
| <b>3. The decisions that matter .....</b>                     | <b>9</b>  |
| <b>4. AI essentials .....</b>                                 | <b>10</b> |
| <b>5. Engineering essentials .....</b>                        | <b>12</b> |
| Backend .....                                                 | 12        |
| Frontend .....                                                | 12        |
| Databases .....                                               | 12        |
| Security .....                                                | 13        |
| Scale, reliability, DevOps .....                              | 13        |
| <b>6. The 35 most likely questions .....</b>                  | <b>14</b> |
| About the project .....                                       | 14        |
| AI / LLM .....                                                | 14        |
| Backend .....                                                 | 15        |
| Frontend .....                                                | 15        |
| Databases .....                                               | 16        |
| Security .....                                                | 16        |
| Tricky follow-ups .....                                       | 16        |
| <b>7. System design in one hour .....</b>                     | <b>17</b> |
| The framework (use it for any design question) .....          | 17        |
| Practice: scale DocMind to 100k users and 10M documents ..... | 17        |
| Two more to sketch on paper .....                             | 18        |
| <b>8. One-page cheat sheet .....</b>                          | <b>19</b> |

# 1. The one-day plan and your pitch

## How to spend the day

| Time                 | Read         | Goal                                                                                         |
|----------------------|--------------|----------------------------------------------------------------------------------------------|
| Morning (2.5 h)      | Sections 1–3 | Say the pitch out loud without notes; draw the architecture from memory; defend 10 decisions |
| Midday (2 h)         | Sections 4–5 | Explain RAG, embeddings, vector search, prompt injection, async, SSE in plain words          |
| Afternoon (2 h)      | Section 6    | Answer the 35 questions aloud; time yourself ( $\leq 1$ minute each)                         |
| Late afternoon (1 h) | Section 7    | Sketch one system design on paper                                                            |
| Evening (30 min)     | Section 8    | Cheat sheet; prepare 3 behavioural stories                                                   |

The full 76-page guide ( [docs/Interview-Study-Guide.pdf](#) ) is your reference if a topic needs more depth.

**Answer shape for every question:** one-sentence answer → one level deeper → an example from DocMind or a trade-off.

## The 60-second pitch

DocMind AI is a full-stack app where you upload PDFs and ask questions about them. Answers stream in token by token and cite the document and page they came from; you can click a citation to see the exact passage.

It's retrieval-augmented generation. On upload, a FastAPI backend extracts text page by page, splits it into overlapping chunks, embeds each chunk with a local model, and stores the vectors in Postgres with pgvector. On a question, I embed it, find the closest chunks *that belong to this user*, and give only those to Claude with strict rules: answer only from these sources, cite them, say "not found" otherwise. If nothing relevant is found, I don't call the LLM at all.

The frontend is Next.js with TypeScript; the stream is Server-Sent Events read with `fetch`, so I can send the auth header and support a Stop button. It's hardened like production — per-user data isolation with tests, upload validation by magic bytes, rate limits, and defences against a planted prompt-injection attack — with 113 automated tests, CI, and an eval set where retrieval finds the right source for 21 of 21 questions.

## The 2-minute extension (if they say "tell me more")


- **Why I built it this way:** answers must be *grounded* (only from your documents), *verifiable* (citations), and *honest* (says "not found").

- **Ingestion:** validate (extension, MIME, %PDF- bytes, 20 MB) → save as a UUID filename → 202 Accepted → background task: pypdf → strip repeated headers → chunk (800 chars, 150 overlap, never across pages) → embed in batches of 32 → insert chunks and set ready in one transaction.
- **Chat:** embed the question → pgvector cosine search filtered by user in SQL → drop scores below 0.45 → prompt with <source id file page> blocks and "treat sources as untrusted data" → stream tokens → send only the citations the answer actually used → save the messages.
- **Lesson learned:** I measured the similarity threshold. Answerable and off-topic questions have *overlapping* scores, so a threshold can't detect unanswerable questions alone; the prompt must refuse. And acronyms ("RPO/RTO") are dense retrieval's blind spot, so hybrid search is my next step.

## Numbers to know

| Fact            | Value                                                                            |
|-----------------|----------------------------------------------------------------------------------|
| Chunks          | 800 characters, 150 overlap, never across pages                                  |
| Embedding model | bge-small-en-v1.5, 384 dimensions, runs locally                                  |
| Retrieval       | cosine distance, HNSW index, top 5, threshold 0.45                               |
| Prompt          | last 6 messages, $\leq 12,000$ characters of sources, answer $\leq 1,024$ tokens |
| Limits          | chat 20/min, upload 10/min, login 5/min, question $\leq 4,000$ chars             |
| Tests           | 100 backend + 13 frontend                                                        |
| Eval            | 22 questions; retrieval 21/21                                                    |


## 2. How DocMind works



**Three tiers:** Next.js frontend → FastAPI backend (all the AI logic) → PostgreSQL + pgvector (rows *and* vectors). The only external call is the LLM.


**Backend layers:** routes (HTTP only) → services (logic) → database. Providers are interfaces ( `EmbeddingProvider`, `LLMProvider` ), so Claude/OpenAI are swapped by config and tests use fakes.

## Upload → ingestion



| Step     | Where                                   | Key point                                                       |
|----------|-----------------------------------------|-----------------------------------------------------------------|
| Validate | <code>utils/upload_validation.py</code> | extension + MIME + %PDF- magic bytes + size, whole batch first  |
| Store    | <code>utils/storage.py</code>           | <uuid>.pdf — user filename never used as a path                 |
| Respond  | <code>routes/documents.py</code>        | 202 Accepted , status=processing , background task scheduled    |
| Parse    | <code>services/pdf_parser.py</code>     | pypdf per page, in a thread (CPU-bound); strip repeated headers |
| Chunk    | <code>services/chunker.py</code>        | paragraph → line → sentence → word; 800/150; per page           |
| Embed    | <code>services/embeddings.py</code>     | batches of 32, in a thread; title + chunk is embedded           |
| Save     | <code>services/ingestion.py</code>      | chunks + ready in one transaction; errors → failed + message    |

## Question → streamed answer



1. **Before streaming:** JWT, validation, rate limit, conversation ownership → normal 4xx errors.
2. Save the user message → `meta` event.
3. **Retrieve** (`services/retriever.py`): `ORDER BY embedding <=> query LIMIT 5`, joined to the user's ready documents, then drop scores < 0.45. Follow-ups also search "previous question + question".
4. **Nothing left?** Send a fixed "not found" answer; **the LLM is not called**.
5. **Prompt** (`services/prompts.py`): system rules + history + numbered `<source>` blocks + question last.
6. **Stream** (`services/llm.py`): each text delta becomes a `token` event.
7. **Citations:** find `[n]` in the answer → send only those → `done`.
8. **Save** the answer. On Stop/disconnect, save the partial answer (in a shielded cancel scope).

**SSE events:** `meta` → `token × N` → `citations` → `done` (or `error` with a code such as `LLM_TIMEOUT`).

## Data model (5 tables)

`users` → `documents` → `chunks` (`embedding vector(384)`, `HNSW index`, `ON DELETE CASCADE`);  
`users` → `conversations` → `messages` (`citations` as JSONB). Every query joins back to `user_id` — that's the security model.

## 3. The decisions that matter

For each: **why this**, the **alternative**, and **when I'd switch**. Say all three and you sound senior.

| Decision                             | Why this                                                                 | Alternative (cost)                                                  | When I'd switch                                         |
|--------------------------------------|--------------------------------------------------------------------------|---------------------------------------------------------------------|---------------------------------------------------------|
| <b>FastAPI (Python)</b>              | Best AI ecosystem; async; validation + docs from type hints              | NestJS (one language, but AI libs are Python-first)                 | TypeScript-only team, AI is a thin API call             |
| <b>pgvector in Postgres</b>          | One datastore: transactions, user filter as a SQL join, one backup       | Pinecone/Qdrant (sync between two stores, more infra)               | 100M+ vectors or vector load must scale separately      |
| <b>Local embeddings (FastEmbed)</b>  | Free, private, no key, fast, small vectors                               | OpenAI/Voyage embeddings (cost, data leaves server)                 | Multilingual or eval shows retrieval is the bottleneck  |
| <b>RAG, not fine-tuning</b>          | Fresh, per-user, citable facts; no retraining                            | Fine-tuning (bad at facts, can't cite, mixes users' data)           | Need a fixed style/format at high volume — combine both |
| <b>Chunks 800/150, per page</b>      | Focused embeddings, boundary overlap, exact page citations               | Fixed windows (cut sentences), semantic chunking (costly)           | Long manuals → larger chunks + parent retrieval         |
| <b>Threshold 0.45 (measured)</b>     | Keeps all answerable questions (lowest 0.508), filters clearly off-topic | 0.35 filters nothing; 0.52 drops a real question                    | A reranker gives calibrated scores to threshold on      |
| <b>SSE via fetch</b>                 | One-way streaming over plain HTTP; auth header; Stop via abort           | WebSockets (stateful, overkill); EventSource (GET only, no headers) | Real-time two-way features                              |
| <b>BackgroundTasks</b>               | No extra infra; 202 + polling                                            | Celery/SQS workers (durable, retries, more ops)                     | Multiple servers, big files, need retries               |
| <b>Provider-agnostic LLM wrapper</b> | Swap Claude/OpenAI by config; one error model                            | LangChain (more abstraction than needed)                            | Many integrations / agent graphs                        |
| <b>Cite only referenced sources</b>  | A chip means "this answer used this passage"                             | Show all retrieved (implies false support)                          | Single provider with native citations                   |
| <b>JWT in localStorage</b>           | Simple for an SPA on a separate API origin                               | httpOnly cookie (no JS access, but needs CSRF protection)           | Production: cookie + refresh tokens                     |
| <b>Rate limits in memory</b>         | One decorator per route; per user, not per IP                            | Redis / API gateway                                                 | More than one backend replica                           |

**The honest summary:** every simplification (background tasks, localStorage token, in-memory limits, local files) has a documented production replacement (queue, cookie, Redis, S3) in `docs/DECISIONS.md`.

## 4. AI essentials

---

**Tokens.** Models read sub-word pieces (~4 characters each in English). You pay per token (output costs more than input), limits are in tokens, and latency grows with output length.

**Context window.** The total tokens the model can see (prompt + answer). Bigger isn't free: cost, latency, and worse use of information in the middle of long prompts. DocMind sends ~5 chunks, 6 history messages,  $\leq 1,024$  answer tokens.

**Temperature.** Low = focused and deterministic, high = varied and error-prone. Use low for factual Q&A. Newer models may replace it with an "effort" setting (DocMind: `effort: low`).

**Embeddings.** Text  $\rightarrow$  vector; similar meaning  $\rightarrow$  close vectors. **Cosine similarity** = angle; on normalized vectors it ranks the same as dot product and L2. BGE models want an instruction prefix on *queries* only.

**RAG in five steps.** Ingest (parse, chunk, embed, store)  $\rightarrow$  retrieve (embed question, nearest neighbours + filters)  $\rightarrow$  augment (prompt with numbered sources)  $\rightarrow$  generate (stream, cite)  $\rightarrow$  post-process (citations, save, log).

**Chunking.** Too big = blurry embeddings and wasted tokens; too small = no context. Recursive splitting on natural boundaries + overlap is the solid default.

**Hybrid search + RRF.** Keyword search finds exact terms (acronyms, codes); vectors find meaning. Merge the two ranked lists with Reciprocal Rank Fusion:  $\text{score} = \Sigma 1 / (60 + \text{rank})$ .

**Reranking.** A cross-encoder reads question and chunk *together* — far more accurate than vector similarity, too slow for everything. Retrieve 20–50, rerank to 5.

**Query rewriting.** Turn "and what about Enterprise?" into a standalone question before retrieval (or DocMind's cheap version: also search with the previous question).

**Hallucination.** Fluent content not supported by sources. Reduce it: skip the LLM when nothing is retrieved, "answer only from sources / say what's missing", require citations, low randomness, evaluate refusals. You reduce and expose hallucinations; you can't eliminate them.

**Prompt engineering.** Clear role, explicit rules, delimiters separating instructions from data, question last, a defined output format, edge cases handled (missing, partial, conflicting info), and a version number — prompts are code.

**Tool calling.** The model asks your code to run a declared function with JSON arguments; *your code* validates, authorizes and runs it. **Agents** are an LLM in a loop of tool calls — use only when steps can't be fixed in advance; add limits and human approval for risky actions.

**Structured outputs.** Enforce a JSON Schema via the API, or validate with Pydantic/Zod and retry.

**Fine-tuning vs RAG vs prompting.** Prompting = behaviour; RAG = knowledge; fine-tuning = style/format/narrow tasks at scale.

**Evaluation.** Measure retrieval (hit-rate@k) and generation (correct facts, faithfulness, citations, correct refusals, injection resistance) separately. LLM-as-judge scales but is biased — validate it against human labels.

**Vector indexes.** **HNSW:** layered graph, great recall, handles inserts, more memory (DocMind's choice). **IVFFlat:** clusters, faster to build, needs data first, lower recall. Both are *approximate* — you trade a little accuracy for a lot of speed.

## 5. Engineering essentials

### Backend

- **Event loop:** one thread juggles many waiting tasks. **Never block it** — CPU work or blocking calls go to `asyncio.to_thread` (DocMind: PDF parsing, embeddings) or to worker processes.
- **Dependency injection:** `Depends(get_current_user)` / `Depends(get_db)` ; tests swap them with `dependency_overrides`.
- **Status codes:** 201 created, **202 accepted (async work)**, 204 deleted, 401 not logged in, 403 not allowed, **404 for other users' data** (don't reveal it exists), 409 conflict, 413 too large, 415 wrong type, 422 invalid, 429 rate limited.
- **SSE vs WebSockets:** SSE = one-way over plain HTTP (streaming answers); WebSockets = two-way and stateful (chat apps, collaboration).
- **Idempotency:** repeating a request has the same effect; use an `Idempotency-Key` for POSTs such as payments.
- **Pagination:** offset (simple, slow on deep pages) vs cursor/keyset (fast, stable).

### Frontend

- **Server vs client components:** server = rendered on the server, no JS shipped; client ( `"use client"` ) = state, effects, events. DocMind's app pages are client components because the token lives in the browser.
- **SSG / SSR / CSR:** build time / per request / in the browser.
- **Keys:** stable IDs, never array indexes for changing lists. **Memo:** `memo` + `useCallback` so only the streaming message re-renders.
- **Stale closure (a real DocMind bug):** a state updater read a variable that was reassigned before it ran → answers stayed blank. Fix: copy into a `const`.
- **Streams:** `res.body.getReader()` + `TextDecoder({stream: true})` + buffer until a blank line. **Stop:** `AbortController`.
- **TypeScript:** strict mode, generics ( `apiFetch<T>` ), discriminated unions for SSE events.

### Databases

- **Indexes:** B-tree (equality/range, foreign keys), GIN (full text, JSONB), HNSW (vectors). Every index slows writes.
- **Transactions (ACID):** DocMind writes chunks + `ready` together. Postgres default isolation = Read Committed. Prefer constraints (unique email) over "check then insert".
- **N+1:** one query per item; fix with eager loading ( `selectinload` ).
- **Pooling:** reuse connections; PgBouncer at scale.
- **Migrations:** versioned (Alembic), expand/contract for zero downtime.

- **SQL vs NoSQL:** relational data with ownership chains and joins → Postgres; flexible documents at huge write scale → MongoDB. **JSONB** for data read as a unit (citations).

## Security

- **AuthN vs AuthZ:** who you are vs what you may do. Most breaches are AuthZ bugs.
- **IDOR:** changing an ID to reach someone else's data → owner check in every query, 404, tests.
- **Passwords:** slow salted hash (bcrypt/Argon2id), never reversible encryption.
- **JWT:** payload is readable; pin the algorithm; check expiry; revocation is hard.
- **CORS** stops *other websites* from reading your API through a user's browser — it is not authentication.
- **CSRF** matters with cookie auth (use SameSite + tokens); DocMind uses a header, so it isn't exposed.
- **XSS:** React escapes text; render Markdown without raw HTML; CSP.
- **Prompt injection (OWASP LLM01):** instructions hidden in data. DocMind's defence: "sources are untrusted data" rule, `<source>` delimiters, escaping fake tags, **no secrets and no tools** — so even a successful injection can't do much. Tested with a planted attack.
- **OWASP LLM Top 10 to name:** prompt injection, sensitive info disclosure, improper output handling, excessive agency, system prompt leakage, vector/embedding weaknesses, misinformation, unbounded consumption.

## Scale, reliability, DevOps

- **Latency:** everything except the LLM is milliseconds; **time to first token** is what users feel. Streaming hides generation time.
- **Scaling out:** stateless API pods + load balancer; move files to S3, limits to Redis, jobs to a queue.
- **Caching:** response, embedding, semantic (similar questions), provider prompt caching.
- **Cost:** skip the LLM when possible, cap tokens, route easy questions to a cheaper model, cache, per-user quotas.
- **Reliability:** timeouts everywhere; retry transient errors with exponential backoff + jitter (only before streaming starts); circuit breaker; fallback provider.
- **Observability:** JSON logs with request IDs; metrics (rate, errors, latency); tracing; AI-specific signals (refusal rate, zero-citation answers, tokens per user).
- **Docker:** multi-stage builds (small, non-root images); Compose for local; CI runs lint, migrations, tests, build.

## 6. The 35 most likely questions

---

Answer aloud, under a minute each.

### About the project

- 1. Walk me through what happens when I ask a question.** Auth and validation → save the question → embed it → pgvector search over *my* chunks → drop weak matches → if nothing is left, a fixed "not found" (no LLM call) → otherwise a prompt with numbered sources → stream tokens over SSE → send the citations the answer used → save the answer.
- 2. What was the hardest part?** Getting "honest" answers: I measured that similarity scores for answerable and off-topic questions overlap, so the threshold only catches clearly off-topic questions and the prompt must handle the rest. Plus making Stop work end to end.
- 3. How do you know it works?** 113 automated tests (security, streaming, refusals, injection) and an eval of 22 labelled questions — retrieval finds the right source for 21 of 21.
- 4. What would you do next?** Hybrid search, a reranker, follow-up question rewriting, a durable ingestion queue, cookie-based auth, OCR.
- 5. What would you change for production?** Queue + workers, S3, Redis for rate limits, httpOnly cookie auth, tracing and dashboards, a secrets manager, cross-provider fallback.

### AI / LLM

- 6. What is RAG, and why not fine-tune?** Retrieve relevant text and answer from it: fresh, private, citable. Fine-tuning teaches style, not reliable facts, can't cite, and would mix users' data.
- 7. What's an embedding? Why cosine?** Meaning as a vector. Cosine compares direction; on normalized vectors it ranks the same as the dot product.
- 8. How did you choose chunk size?** ~800 characters (~200 tokens) + 150 overlap, split on natural boundaries, never across pages. Focused, enough context, exact page citations.
- 9. How do you stop hallucinations?** Don't call the LLM without context; "only from sources, say what's missing"; mandatory citations shown to the user; low randomness; refusal questions in the eval.
- 10. How did you pick the similarity threshold?** Measured: answerable 0.508–0.84, off-topic 0.416–0.541. 0.45 keeps every real question.
- 11. What's hybrid search?** Keyword (exact terms) + vector (meaning), merged by rank with RRF. It fixes acronyms like "RPO/RTO", my weakest eval question.
- 12. What does a reranker do?** A cross-encoder scores (question, chunk) pairs together — more accurate; run it on the top 20–50 candidates.

**13. How do you evaluate a RAG app?** Retrieval (hit-rate@k) and answers (facts, faithfulness, citations, refusals, injection) separately, on a versioned set, after every change. LLM-as-judge for paraphrases, validated against humans.

**14. What is prompt injection, and how did you defend?** Instructions hidden in data. Untrusted-data rule, `<source>` delimiters, escaped fake tags, and least privilege: no secrets in the prompt, no tools, per-user data, output never executed. Tested with a planted attack.

**15. Tool calling and agents?** The model requests a function call; my code validates, authorizes and runs it. Agents loop over this — only when steps can't be fixed in advance, with limits and approvals.

**16. How do you control LLM cost?** Skip calls when possible, cap context and output tokens, cache, route to cheaper models, per-user quotas, and log tokens per request.

## Backend

**17. Why FastAPI?** Async, validation and docs from type hints, and Python's AI ecosystem.

**18. What if you call blocking code in `async def`?** It freezes every request on that worker. Use `to_thread` or worker processes.

**19. Why 202 for uploads?** Accepted; processing continues in the background; the client polls for `ready / failed`.

**20. SSE or WebSockets?** SSE: one-way over plain HTTP, ideal for token streaming. WebSockets: two-way, stateful.

**21. BackgroundTasks vs a queue?** In-process and simple vs durable, retried and independently scaled. I'd switch at scale.

**22. What happens when the user presses Stop?** The fetch aborts → the server cancels the stream → the provider stream closes (no more tokens billed) → the partial answer is saved in a shielded scope.

## Frontend

**23. Server vs client components?** Server: no JS shipped, fetches on the server. Client: interactivity. My pages are client components because the token is in the browser.

**24. How do you read a stream in the browser?** `getReader() + TextDecoder({stream: true})` + buffer events until a blank line; `AbortController` to cancel.

**25. How do you avoid re-renders while streaming?** Memoize message components, stable callbacks, keep state local.

**26. Tell me about a bug you found.** A stale closure: a state updater read a variable reassigned before React ran it, so answers stayed blank. An end-to-end test caught it; fix: capture into a `const`.

## Databases

**27. Why pgvector and not Pinecone?** One datastore, transactions, user filters as joins, no syncing. Switch at hundreds of millions of vectors.

**28. HNSW vs IVFFlat?** Graph vs clusters. HNSW has better recall and handles inserts; IVFFlat builds faster but needs data first.

**29. What's N+1?** One query per list item; fix with eager loading.

**30. Isolation levels?** Read Committed (default), Repeatable Read, Serializable. Many races are better solved with a unique constraint.

## Security

**31. What's IDOR, and how did you prevent it?** Changing an ID to reach others' data → owner check inside every query, 404, UUIDs, tests for each resource.

**32. How do you store passwords?** bcrypt (cost 12), salted, slow; Argon2id is the modern first choice.

**33. JWT pitfalls?** Readable payload, pin the algorithm (`alg: none` attack), check expiry, revocation is hard, storage trade-offs.

**34. CORS, CSRF, XSS in one breath?** CORS = other sites can't read my API via your browser; CSRF = forged cookie requests (not an issue with header tokens); XSS = injected script (React escaping, no raw HTML, CSP).

## Tricky follow-ups


**35. Quick answers:** - *Scanned PDF?* No text layer → marked `failed` today; add OCR (Tesseract or a cloud OCR service) with page numbers. - *500-page document?* Queue worker, page batches with progress, larger embedding batches, section-aware chunking; retrieval is still fine (~1,500 chunks). - *Multi-language?* Multilingual embedding model + re-embed, answer in the user's language, language-aware keyword search, non-English eval questions. - *PII?* Minimize, per-user access, never log contents, encrypt, check provider data retention, deletion that cascades, optional PII redaction before the LLM. - *Conflicting sources?* The prompt says who says what and prefers authoritative documents; later, rank by date/authority metadata.

# 7. System design in one hour

## The framework (use it for any design question)

1. **Requirements:** what it does + scale, latency, availability, security, cost. Ask questions.
2. **Estimates:** users, requests/second, data size.
3. **API + data model:** the few key endpoints and tables.
4. **Boxes and arrows:** walk one request through.
5. **Deep dives:** the 2–3 hardest parts.
6. **Bottlenecks, failures, trade-offs, monitoring.**

## Practice: scale DocMind to 100k users and 10M documents



**Estimates.** 10M docs × 20 pages × 3 chunks ≈ **600M chunks** → 384 × 4 bytes each ≈ **~920 GB of vectors** (half that with 16-bit `halfvec`). Too big for one Postgres node. Chat load is modest (~7 questions/s at peak); LLM cost and time to first token matter more.

### Design.

- CDN for the frontend; **stateless API pods** behind a load balancer.
- **Uploads straight to S3** (presigned URLs) → **queue** → **autoscaled ingestion workers** → **embedding service** (GPU batches or a hosted API).
- **Vectors:** pgvector **partitioned by tenant** (with 16-bit vectors) *or* a dedicated vector DB (Qdrant/Pinecone) kept in sync via an outbox table.
- Postgres primary + **read replicas**; PgBouncer; **Redis** for rate limits and caches.

- **Better retrieval:** hybrid search + reranker.
- **LLM layer:** circuit breaker, fallback provider, prompt caching, model routing, per-tenant token quotas.
- **Observability:** traces, time to first token, refusal rate, tokens per tenant, queue depth.

**Bottlenecks → fixes.** Vector memory → partitioning/quantization/vector DB. Bulk imports → GPU workers, per-tenant fairness. Filtered search recall for small tenants → partition or exact search. Provider limits → quotas and fallback. Deletion at scale → delete across S3, vectors and rows, with an audit trail.

## Two more to sketch on paper

**AI support chatbot with tools.** RAG over help articles + narrow allow-listed tools ( `get_order`, `track_shipment`, `create_return`, `issue_refund` ). Validate arguments, **authorize against the logged-in customer** (never trust the model), confirm or hand to a human for risky actions, limit steps per turn, audit log, human handoff with a summary.

**Multi-tenant SaaS.** Pool (shared tables + `tenant_id` + **Postgres Row-Level Security**) by default; schema-per-tenant or **database-per-tenant** for enterprise customers. Tenant-scoped vector filters and cache keys, per-tenant keys (KMS), quotas and cross-tenant tests.

## 8. One-page cheat sheet

**Pitch:** Upload PDFs, ask questions, get streamed answers that cite document + page, and an honest "not found" when the documents don't say. RAG with FastAPI + pgvector + local embeddings + Claude; Next.js frontend; hardened and tested.

| Pipeline | Numbers                                                                                               |
|----------|-------------------------------------------------------------------------------------------------------|
| Parse    | pypdf page by page; strip running headers; scanned → failed                                           |
| Chunk    | 800 chars, 150 overlap, recursive (¶ → line → sentence → word), never across pages                    |
| Embed    | bge-small-en-v1.5, 384-d, normalized, batch 32, local ONNX; query prefix instruction                  |
| Store    | Postgres 16 + pgvector, HNSW vector_cosine_ops, one transaction with status=ready                     |
| Retrieve | <=> cosine distance, user filter in SQL, TOP_K 5, MIN_SIMILARITY 0.45                                 |
| Prompt   | system rules + last 6 messages + <source id file page> blocks (≤ 12k chars) + question last           |
| Generate | claude-opus-5-5, effort low, max_tokens 1024, 60 s timeout, 2 retries before stream, refusal fallback |
| Stream   | SSE: meta → token... → citations → done / error; fetch + ReadableStream + AbortController             |
| Cite     | parse [n], show only cited sources, chip opens /api/chunks/{id}                                       |

**Limits:** chat 20/min/user · upload 10/min/user; 20 MB, 10 files · login 5/min/IP · question ≤ 4,000 chars · JWT 60 min · bcrypt cost 12.

**Quality:** 113 tests (100 backend, 13 frontend) · eval 22 questions · retrieval hit-rate 21/21 · threshold measured: answerable 0.508–0.84 vs off-topic 0.416–0.541 (overlap).

### Key terms (one line each):

- **RAG** — retrieve relevant text, then generate from it. **Grounding** — answer only from provided sources.
- **Embedding** — meaning as a vector. **Cosine** — angle-based similarity; = dot product on unit vectors.
- **HNSW** — layered graph ANN index; m, ef\_construction, ef\_search. **IVFFlat** — cluster buckets; lists, probes.
- **Hybrid + RRF** — keyword + vector; merge by  $1/(60 + \text{rank})$ . **Reranker** — cross-encoder on top-k.
- **Prompt injection** — instructions hidden in data; defend with delimiters, untrusted-data rule, least privilege.
- **IDOR** — access by changing an ID; fix with ownership in every query, 404, tests.
- **SSE** — one-way HTTP stream; **WebSocket** — two-way stateful connection.
- **202 Accepted** — async processing; **idempotency key** — safe retries of POST.
- **Event loop** — never block it; `asyncio.to_thread` for CPU/blocking work.
- **N+1** — fix with eager loading. **ACID**; Postgres default isolation = Read Committed.
- **TTFT** — time to first token, the latency that users feel with streaming.

### Talking points (have a story for each):

1. Measured the threshold instead of guessing — and found that a threshold can't detect unanswerable-but-on-topic questions.
2. Prompt injection is contained by design: no secrets, no tools, per-user data, output never executed — plus tests and an eval question.
3. Stop button end to end: AbortController → server cancellation → provider stream closed → partial answer saved in a shielded scope.
4. Stale-closure bug caught by an end-to-end test, not unit tests.
5. Every simplification has a documented production path: BackgroundTasks → queue, localStorage → httpOnly cookie, in-memory limits → Redis, local files → S3.

**Next steps:** answer eval with live model → hybrid search → reranker → query rewriting → durable queue → cookie auth → OCR → feedback + semantic cache.

**Your questions for them:** How do you evaluate AI features? Path from prototype to production? Biggest technical challenge this year?