

Q2 2026 Engineering Report

Period: 1 April – 30 June 2026 · Prepared by the VP of Engineering

1. Executive summary

Q2 was a strong quarter for reliability and delivery speed. Platform uptime reached 99.93% against a target of 99.9%, p95 API latency fell from 340 ms to 210 ms after the move to regional read replicas, and cloud spend dropped 11% quarter over quarter thanks to automatic storage tiering. The main setback was one SEV1 incident on 14 May.

2. Key metrics

Metric	Q1 2026	Q2 2026	Target
Platform uptime	99.87%	99.93%	99.9%
p95 API latency	340 ms	210 ms	< 250 ms
Deployments per week (avg)	11	18	15
Lead time for changes	2.9 days	1.6 days	< 2 days
Change failure rate	6.8%	4.2%	< 5%
Mean time to recovery (MTTR)	71 min	38 min	< 60 min
Cloud cost (quarter)	USD 463,000	USD 412,000	USD 430,000
Engineering headcount	55	64	66

Deployment frequency, lead time, change failure rate and MTTR are the four DORA metrics, which we use as our main delivery health indicators.

3. Incidents

3.1 SEV1: Sync outage on 14 May 2026

Between 09:12 and 09:59 PKT the desktop and mobile sync service was unavailable for all customers, a total of 47 minutes. The root cause was an expired TLS certificate on an internal load balancer that was not covered by automated certificate renewal. No customer data was lost.

Follow-up actions: all internal certificates were moved to automated renewal, a monitor now alerts 30 days before any certificate expiry, and certificate inventory was added to the quarterly access review.

3.2 SEV2 incidents

There were three SEV2 incidents in the quarter: delayed search indexing on 3 April (2 hours 10 minutes), elevated upload errors in the EU region on 22 May (35 minutes), and a failed database migration on 18 June that was rolled back within 25 minutes.

4. Team

Engineering grew from 55 to 64 people, with nine new hires: four backend engineers, two frontend engineers, one ML engineer, one SRE and one engineering manager. Voluntary attrition for the quarter was 2.1%. The new Platform AI team was formed in May and is led by Sana Qureshi.

5. Delivered in Q2

- Regional read replicas for the metadata database, which drove the latency improvement.
- Automatic storage tiering that moves files not accessed for 180 days to cheaper cold storage.
- Vault release 7.3 with passwordless login via passkeys.
- Migration of CI pipelines to ephemeral runners, cutting average build time from 14 minutes to 6 minutes.

6. Q3 2026 roadmap

Initiative	Owner	Target date
Vault AI Search (semantic search across documents using embeddings)	Platform AI team	15 September 2026
Middle East region (Bahrain) general availability	Infrastructure	30 August 2026
Offline mode for the Android app	Mobile team	30 September 2026
Reduce p95 API latency to 180 ms	Core API team	30 September 2026

The largest risk for Q3 is the AI Search launch: it depends on the new vector index service, and the Platform AI team is still hiring two engineers. Customer content used by AI Search will not be used for model training, in line with the Security and Data Protection Policy.